


LORA: LOW-RANK ADAPTATION OF LARGE LANGUAGE MODELS

Edward J. Hu and Yelong Shen and Phillip Wallis and Zeyuan Allen-Zhu and Yuanzhi Li and Shean Wang
and Lu Wang and Weizhu Chen

Presented by,
Anagha M B, Gokul P V, and Chandu Dadi

personalized models



Challenges with Existing Adaptation Methods

Adapter Layers:

Small modules inserted into transformer layers. Their purpose is to modify the model's behavior for specific tasks without fine-tuning the entire network.

Advantages: Parameter-efficient; supports multi-task learning

Problems:

- **Inference Latency:** Additional depth increases sequential computations.
- **Scalability Issues:** Requires GPU synchronization (e.g., AllReduce) in large-scale deployments.

Prompt Tuning:

Modifies **input tokens** with task-specific prefixes.

Example:

Input sequence: "Translate English to French: What is your name?"

With prompt tuning: "[TASK_SPECIFIC] Translate English to French: What is your name?"

Advantages: Keeps model weights frozen; lightweight.

Problems:

- **Optimization Difficulty:** Training prompt embeddings is unstable.
- **Sequence Length Reduction:** Prefix tokens reduce space for task input.

Eckart-Young Theorem

Given a matrix $A \in \mathbb{R}^{m \times n}$ and its singular value decomposition:

$$A = U \Sigma V^T$$

where $\Sigma = \text{diag}(\sigma_1, \sigma_2, \dots, \sigma_r)$, with $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$ (singular values), the best rank- k approximation of A , in the Frobenius norm, is:

$$A_k = U_k \Sigma_k V_k^T$$

where Σ_k is the diagonal matrix containing the top k singular values, and U_k , V_k are the corresponding singular vectors.

Key Ideas of LoRA

- Instead of updating all model parameters, LoRA injects trainable low-rank matrices into the pretrained weights.
- Pretrained weights remain frozen, reducing computational overhead.
- Parameter updates are represented as the product of two low-rank matrices.
- Gradients scaled by α/r




Figure 1: Our reparametrization. We only train A and B .

$$h = W_0 x + \Delta W x = W_0 x + B A x$$

LoRA Memory Efficiency

- The low-rank structure significantly reduces memory overhead and computational complexity during training.

Storage: $d \times k \rightarrow r(d+k)$

- Example: For $d=512$, $k=512$, $r=8$
 - Original: 262,144 parameters
 - LoRA: 8,192 parameters (96.9% reduction)
- Maintains expressiveness of full-rank updates with far fewer trainable parameters.



https://colab.research.google.com/drive/1DzArDRxgcPiu32fwWd3BgREy9SKRcOJ_



<https://colab.research.google.com/drive/142Wy4B5cVUWnzS2JaSt9jPsyVcYTtB3O>

Applications of LoRA

1. Natural Language Processing (NLP)

- a. **Text Classification:** Fine-tuning LLMs like RoBERTa, BERT, or GPT for sentiment analysis, spam detection, or topic classification with minimal parameter updates.
- b. **Question Answering (QA):** Tuning models like RoBERTa or GPT for domain-specific Q&A applications, such as customer support or medical FAQs.

2. Personalized AI and Chatbots

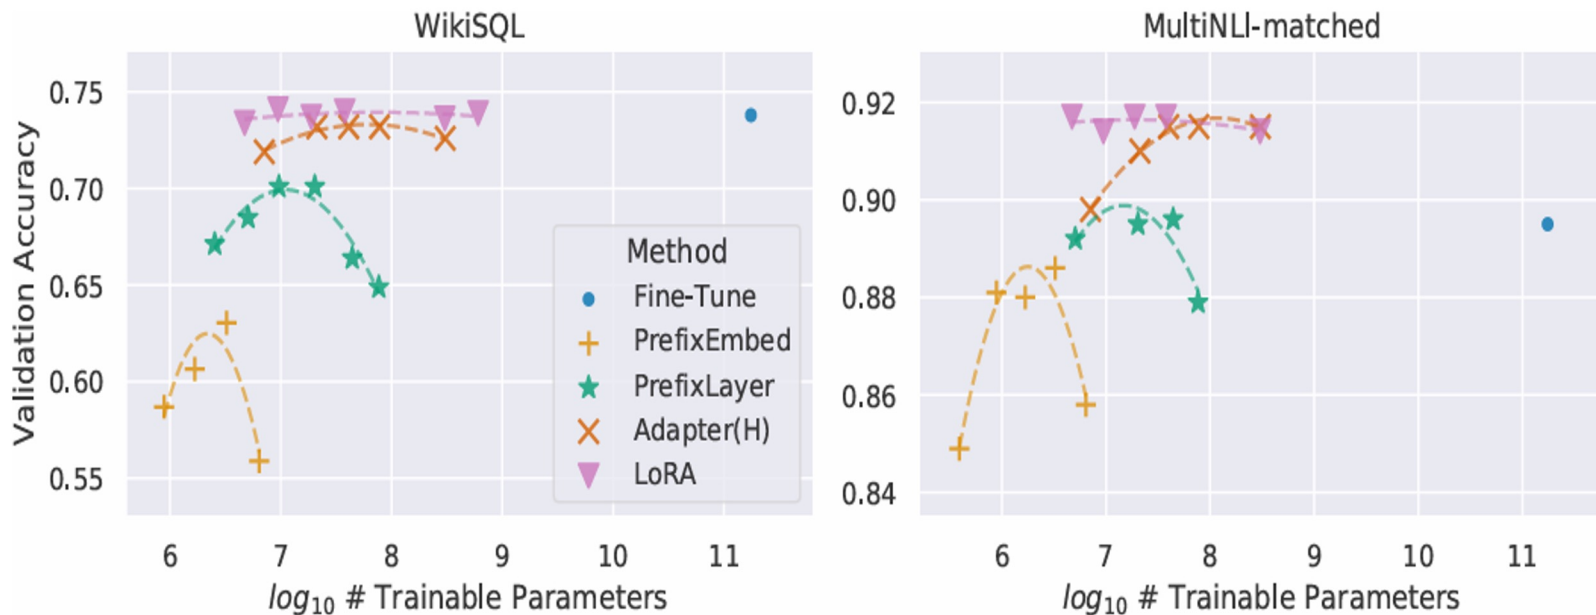
3. Generative AI

EMPIRICAL EXPERIMENTS

Model & Method	# Trainable Parameters	MNLI	SST-2	MRPC	CoLA	QNLI	QQP	RTE	STS-B	Avg.
RoB _{base} (FT)*	125.0M	87.6	94.8	90.2	63.6	92.8	91.9	78.7	91.2	86.4
RoB _{base} (BitFit)*	0.1M	84.7	93.7	92.7	62.0	91.8	84.0	81.5	90.8	85.2
RoB _{base} (Adpt ^D)*	0.3M	87.1 \pm .0	94.2 \pm .1	88.5 \pm 1.1	60.8 \pm .4	93.1 \pm .1	90.2 \pm .0	71.5 \pm 2.7	89.7 \pm .3	84.4
RoB _{base} (Adpt ^D)*	0.9M	87.3 \pm .1	94.7 \pm .3	88.4 \pm .1	62.6 \pm .9	93.0 \pm .2	90.6 \pm .0	75.9 \pm 2.2	90.3 \pm .1	85.4
RoB _{base} (LoRA)	0.3M	87.5 \pm .3	95.1\pm.2	89.7 \pm .7	63.4 \pm 1.2	93.3\pm.3	90.8 \pm .1	86.6\pm.7	91.5\pm.2	87.2
RoB _{large} (FT)*	355.0M	90.2	96.4	90.9	68.0	94.7	92.2	86.6	92.4	88.9
RoB _{large} (LoRA)	0.8M	90.6\pm.2	96.2 \pm .5	90.9\pm1.2	68.2\pm1.9	94.9\pm.3	91.6 \pm .1	87.4\pm2.5	92.6\pm.2	89.0
RoB _{large} (Adpt ^P)†	3.0M	90.2 \pm .3	96.1 \pm .3	90.2 \pm .7	68.3\pm1.0	94.8\pm.2	91.9\pm.1	83.8 \pm 2.9	92.1 \pm .7	88.4
RoB _{large} (Adpt ^P)†	0.8M	90.5\pm.3	96.6\pm.2	89.7 \pm 1.2	67.8 \pm 2.5	94.8\pm.3	91.7 \pm .2	80.1 \pm 2.9	91.9 \pm .4	87.9
RoB _{large} (Adpt ^H)†	6.0M	89.9 \pm .5	96.2 \pm .3	88.7 \pm 2.9	66.5 \pm 4.4	94.7 \pm .2	92.1 \pm .1	83.4 \pm 1.1	91.0 \pm 1.7	87.8
RoB _{large} (Adpt ^H)†	0.8M	90.3 \pm .3	96.3 \pm .5	87.7 \pm 1.7	66.3 \pm 2.0	94.7 \pm .2	91.5 \pm .1	72.9 \pm 2.9	91.5 \pm .5	86.4
RoB _{large} (LoRA)†	0.8M	90.6\pm.2	96.2 \pm .5	90.2\pm1.0	68.2 \pm 1.9	94.8\pm.3	91.6 \pm .2	85.2\pm1.1	92.3\pm.5	88.6
DeB _{XXL} (FT)*	1500.0M	91.8	97.2	92.0	72.0	96.0	92.7	93.9	92.9	91.1
DeB _{XXL} (LoRA)	4.7M	91.9\pm.2	96.9 \pm .2	92.6\pm.6	72.4\pm1.1	96.0\pm.1	92.9\pm.1	94.9\pm.4	93.0\pm.2	91.3

EMPIRICAL EXPERIMENTS

GPT-3 175 B



CONCLUSION

Fine-tuning enormous language models is expensive in terms of the hardware required and the storage/switching cost for hosting independent instances for different tasks.

LoRA, an efficient adaptation strategy that neither introduces inference latency nor reduces input sequence length while retaining high model quality.

References

- <https://arxiv.org/pdf/2106.09685>
- <https://magazine.sebastianraschka.com/p/practical-tips-for-finetuning-llms>
- <https://towardsdatascience.com/understanding-lora-low-rank-adaptation-for-finetuning-large-models-936bce1a07c6>
- <https://github.com/huggingface/smol-course>

Thank you!